


Hate-speech classification recall by model and condition



* $p_{\text{FDR}} < 0.05$, ** $p_{\text{FDR}} < 0.01$, *** $p_{\text{FDR}} < 0.001$ (Wilcoxon signed-rank vs. zero-shot, BH-corrected across models)